

PATHspider II: The Tutorial

Iain R. Learmonth

University of Aberdeen

May 16, 2018

Active Measurement of Path Transparency

- Methodology:
 - 1 Throw packets at the Internet
 - 2 See what happens
- Ideal: two-ended A/B testing
- Scalable: one-ended A/B testing
- Multiple sources: **isolate** on-path from near-target impairment




Figure: One-Ended vs. Two-Ended Testing

History

ecns spider

- The original implementation supported by mPlane/RITE
- Three distinct components:
 - DNS List Resolver
 - QoF Flow Meter
 - Active Traffic Generator
- Used hardcoded `sysctl(1)` and `iptables(1)` commands to cause packets to be emitted with various ECN-related flags
- Source code:
<https://github.com/britram/ecns spider>




Figure: Original Architecture

History

History ecnsipider

- The original implementation supported by mPlane/RITE
- Three distinct components:
 - DNS List Resolver
 - QoF Flow Meter
 - Active Traffic Generator
- Used hardcoded `sysctl(1)` and `iptables(1)` commands to cause packets to be emitted with various ECN-related flags
- Source code:
<https://github.com/britram/ecnsipider>



Figure: Original Architecture

QoF is an IPFIX Metering and Exporting process derived from the YAF flowmeter, designed for passive measurement of per-flow performance characteristics.

While it was fast, it was only able to export flow properties it already knew about, and could not be easily extended.

If you are interested in network measurement, you may still like to read about Internet Protocol Flow Information Export (IPFIX) [3]. It was created based on the need for a common, universal standard of export for Internet Protocol flow information from routers, probes and other devices that are used by mediation systems, accounting/billing systems and network management systems to facilitate services such as measurement, accounting and billing.

History

ecns spider Results

- ECN negotiation was found to be successful for 56.17% of hosts connecting for IPv4, 65.41% for IPv6, from the Alexa top 1 million list [8]
- This continues a trend ETH started observing with *ecns spider* in 2013 [6]




Figure: ECN Support in the Alexa Top 1 Million

PATHspider 1.0¹

- Architecture based closely on the original ecnspider
- Generalised to support more than just ECN
 - Added TCP Fast Open and DiffServ Codepoints
- Still performing A/B testing, but with more A/B tests
- Replaced QoF with a Python flowmeter implementation using python-libtrace
- Began to develop a generalised measurement methodology for path transparency testing
- Published at 2016 Applied Networking Research Workshop [7]

Plugin Architecture

- The plugin architecture was not as generalised as it could have been
- Plugin methods:
 - `config_zero`
 - `config_one`
 - `connect`
 - `post_connect`
 - `create_observer`
 - `merge`

Built-In Flowmeter

- PATHspider's built in flow meter is extensible via the plugin architecture
- Using python-libtrace to dissect packets, any flow property imaginable can be reported back based on the raw packets:
 - ECN negotiation (IP/TCP headers)
 - Bleaching of bits, dropping of options
 - Checksum recalculations

PATHspider 1.0 Results


We presented some initial findings along with the publication of PATHspider 1.0 [7]:

Explicit Congestion Notification (ECN)

State of ECN server-side deployment, as measured from a Digital Ocean vantage point in Amsterdam on 13th June 2016:

	IPv4	IPv6	all
No ECN connectivity issues	99.5%	99.9%	99.5%
ECN successfully negotiated	70.0%	82.8%	70.5%

ECN negotiation by Alexa rank bin:



DiffServ Code Points (DSCP)

Initial study: 10,006 of 96,978 (10.31%) of Alexa Top 100k

websites had unexpected, non-zero DSCP values. More

measurement was needed to better characterize these anomalies.

TCP Fast Open (TFO)

Initial study: 330 IPv4 and 32 IPv6 addresses of Alexa Top 1M

are TFO-capable (of which 278 and 28 respectively are Google

properties). DDoS prevention services, enterprise firewalls, and

CPE tend to interfere with TFO. More measurement was

necessary to analyze impairments.

PATHspider 1.0 Results

Higher ranked servers tended to disable (or not support to begin with) ECN. This is likely due to the specialised nature of services that have to handle such large volumes of traffic. They may be using entirely custom codebases, or are otherwise tuned and optimised.

For a more comprehensive measurement study on the use and impairments to use of DiffServ Codepoints in the Internet, see [4].

PATHspider 1.0 Results


We presented some initial findings along with the publication of PATHspider 1.0 [7].

Explicit Congestion Notification (ECN)

State of ECN server side deployment, as measured from a Digital Ocean vantage point in Amsterdam on 19th June 2018.

	IPV4	IPV6	all
No ECN immediately before	101/111	101/100	101/105
ECN immediately before	10/10	12/10	11/10

ECN negotiation by client rank 5th.



DiffServ Code Points (DSCP)

Initial study: 100/100 of 100/100 (100/100) of down. Top 10th percentile had approximately 100/100 DSCP values. More measurement was needed to better characterise these anomalies.

TCP Fast Open (TFO)

Initial study: 100/100 and 10/100 of down. Top 10th percentile had approximately 100/100 and 10/100 of down. More measurement was needed to better characterise these anomalies.

PATHspider 2.0

- Architecture changed to add a flow combiner
- Generalised to support more than just A/B testing
 - Any permutation of any number of tests
- Replaced PATHspider's HTTP code with cURL
- Added framework for packet forging based plugins using Scapy
- Completely rewritten (in Go) target list resolver
- Observer modules usable for standalone passive observation or analysis
- Source code: <https://github.com/mami-project/pathspider/tree/2.0.0/>

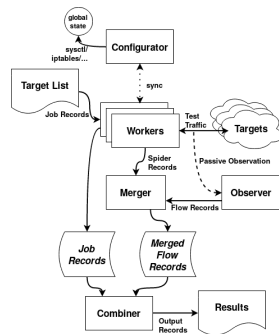


Figure: New Architecture

PATHspider II: The Tutorial

└─PATHspider 2.0

PATHspider 2.0

- Architecture changed to add a flow combiner
- Generalised to support more than just A/B testing
 - Any permutation of any number of tests
- ◆ Replaced PATHspider's HTTP code with cURL
- ◆ Added framework for packet forging based plugins using Scapy
- ◆ Completely rewritten (in Go) target list resolver
- ◆ Observer modules usable for standalone passive observation or analysis
- ◆ Source code: <https://github.com/mami-project/pathspider/tree/2.0.0/>



Figure: New Architecture

The combiner thread holds a table of merged flows and waits for $|flows| = |jobs|$. Conditions are generated based on the combined flows.

Plugin Types

- Synchronised (traditional ecnspider)
 - ECN, DSCP
- Desynchronised (traditional ecnspider, no configurator)
 - TFO, H2, TLS NPN/ALPN
- Forge (new in PATHspider 2.0!)
 - Evil Bit, UDP Zero Checksum, UDP Options
- Single (new, and fast)
 - Various TCP Options

└ Plugin Types

Plugin Types

- ◆ Synchronised (traditional ecnspider)
 - ECN, DSCP
- ◆ Desynchronised (traditional ecnspider, no configurator)
 - TFO, H2, TLS NP/N, ALPN
- ◆ Forge (new in PATHspider 2.0!)
 - Eol Bit, UDP Zero Checksum, UDP Options
- ◆ Single (new, and fast)
 - Various TCP Options

The desynchronized plugins will run more quickly than synchronized plugins while still using the real network stack. Forge plugins will run slower as there is overhead in the Scapy packet generation that doesn't exist in optimised kernel stacks.

Connection Helpers

- Instead of writing client code, use the code that already exists
- In the `pathspider.helpers` module:
 - DNS (`dnslib`)
 - HTTP/HTTPS (`pycURL`)
 - TCP (Python `socket`)
- For synchronised plugins, just use the helper
- For desynchronised plugins, the helpers are customisable, e.g. `cURL` helpers accept arbitrary `CURL_OPTS`

└ Connection Helpers

Connection Helpers

- Instead of writing client code, use the code that already exists
- In the `pathspider.helpers` module:
 - DNS (`dnlib`)
 - HTTP/HTTPS (`pycURL`)
 - TCP (Python `socket`)
- For synchronised plugins, just use the helper
- For desynchronised plugins, the helpers are customisable, e.g. `cURL` helpers accept arbitrary `CURLOPTs`

todo: Pointer to the `pycURL` documentation, and talk about upstreaming `CURLOPTs`

Synchronized Plugin

todo: Write this

Desynchronized Plugin

todo: Write this

Forge Plugin

todo: Write this

Single Plugin

todo: Write this

Observer Modules

- While these used to be part of plugins in PATHspider 1.0, they are now independent and so can be reused across multiple plugins:
 - BasicChain, DNSChain, DSCPChain, ECNChain, EvilChain, ICMPChain, TCPChain, TFOChain
- These can also be used together, limiting each chain to just a single layer and letting the combiner produce conditions
- Chains can produce information to be consumed by other chains later in the list
- These can be used independently of a PATHspider measurement:

```
1 irl@z~$ pspdr observe tcp ecn
```

Listing 1: Running the PATHspider Observer independently

Target List Resolution

- Hellfire is a parallelised DNS resolver. It is written in Go and for the purpose of generating input lists to PATHspider, though may be useful for other applications
- Can use many sources for inputs:
 - Alexa Top 1 Million Global Sites
 - Cisco Umbrella 1 Million
 - Citizen Lab Test Lists
 - OpenDNS Public Domain Lists
 - Comma-Separated Values Files
 - Plain Text Domain Lists
- More on this later



Packet Forging

- PATHspider uses the Scapy library for Python for packet forging
- This is the most flexible method of creating new measurement plugins for PATHspider

Make a Packet

- Scapy packets are constructed layer by layer
- While you can specify raw bytes, Scapy provides a number of useful classes for common protocols, which makes things a lot easier

Make a Packet

- Scapy must be launched with `sudo` as we will need to use “raw” sockets to emit forged packets.
- Note also the command is `scapy3`, to ensure we are running the Python 3 version.

```
1 irl@z:~$ sudo scapy3
3
5      aSPY//YASa
6      apyyyyCY////////YCa
7      sY////////YSpcs  scpCY//Pp
8  ayp ayyyyyyySCP//Pp      syY//C
9  AYAsAYYYYYYYY//Ps      cY//S
10     pCCCY//p      cSSps y//Y
11     SPPPP//a      pP//AC//Y
12     A//A      cyP////C
13     p///Ac      sC///a
14     P////YCpc      A//A
15     sccccp///pSP///p      p//Y
16     sY////////y caa      S//P
17     cayCyayP//Ya      pY/Ya
18     sY/PsY////YCc      aC//Yp
19     sc  sccaCY//PCypaapyCP//YSs
20         spCPY////////YPSps
21         ccaacs
22
23                                     using IPython 5.5.0
24
25  >>>
```

Welcome to Scapy
Version 2.4.0

<https://github.com/secdev/scapy>

Have fun!

We are in France, we say Skappee.
OK? Merci.

— Sebastien Chabal

Listing 2: Launching Scapy

Make a Packet

IPv4 Header - Create and Dissect

```
1 >>> IP()
  <IP  |>
3 >>> i = IP()
  >>> i.summary()
5 '127.0.0.1 > 127.0.0.1 hopopt'
  >>> i.display()
7 ###[ IP ]###
  version= 4
  ihl= None
  tos= 0x0
11  len= None
  id= 1
13  flags=
  frag= 0
15  ttl= 64
  proto= hopopt
17  chksum= None
  src= 127.0.0.1
19  dst= 127.0.0.1
  \options\
```

Listing 3: Creating and Dissecting an IPv4 Header

Make a Packet

IPv4 Header - Customize

```
>>> i = IP(src="192.0.2.1",dst="198.51.100.1",ttl=10)
>>> i
<IP ttl=10 src=192.0.2.1 dst=198.51.100.1 |>
>>> i.summary()
'192.0.2.1 > 198.51.100.1 hopopt'
>>> i.display()
###[ IP ]###
version= 4
ihl= None
tos= 0x0
len= None
id= 1
flags=
frag= 0
ttl= 10
proto= hopopt
chksum= None
src= 192.0.2.1
dst= 198.51.100.1
\options\
```

Listing 4: Customizing an IPv4 Header²

Make a Packet

IPv4 Header - Bonus: Export a Dissection

```
>>> i.pdfdump()
```

Listing 5: Create a PDF Export of a Dissection of the IP Header




Figure: PDF Export of IP Header Dissection

Make a Packet

TCP Header: Create and Dissect

```
1 >>> TCP()
  <TCP  |>
3 >>> t = TCP()
  >>> t.summary()
5 'TCP ftp_data > http S'
  >>> t.display()
7 ###[ TCP ]###
  sport= ftp_data
  dport= http
  seq= 0
  ack= 0
  dataofs= None
  reserved= 0
  flags= S
  window= 8192
  chksum= None
  urgptr= 0
  options= []
```

Listing 6: Creating and Dissecting a TCP Header

Make a Packet

TCP Header: Customizing

```
1 >>> t = TCP(dport=443)
2 >>> t
3 <TCP dport=https |>
4 >>> t.summary()
5 'TCP ftp_data > https S'
6 >>> t.display()
7 ##### TCP #####
8 sport= ftp_data
9 dport= https
10 seq= 0
11 ack= 0
12 dataofs= None
13 reserved= 0
14 flags= S
15 window= 8192
16 chksum= None
17 urgptr= 0
   options= []
```

Listing 7: Customizing a TCP Header

Modifying a Packet

todo: Modify a packet

Make a Packet

Sticking the Pieces Together

- The / operator is used to join layers together.
- Scapy will automatically set fields, such as the IP Protocol field, when you do this.
- When dissecting, Scapy will automatically choose the dissector to use based on fields such as the IP Protocol field.

```
>>> p=i/t
>>> p.summary()
'IP / TCP 192.0.2.1:ftp_data > 198.51.100.1:https S'
>>> p.display()
[... output snipped ...]
```

Listing 8: Sticking the IP and TCP Headers Together

Make a Packet

View in Wireshark

```
>>> wrpcap("/tmp/scapy.pcap", [p])
```

Listing 9: Exporting a PCAP File from Scapy

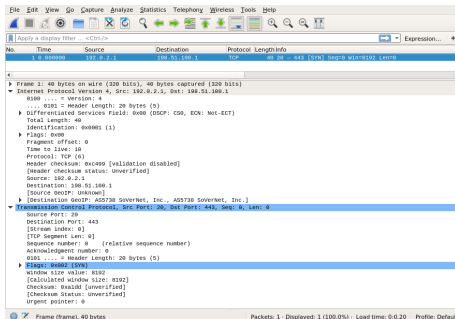


Figure: Dissection of the packet created in Scapy, in Wireshark

Send a Packet

- The `sr1()` function sends a single packet, and returns a single packet if a reply is received.
- Start Wireshark capturing before executing the `sr1()` function.

```
1 >>> p=IP(dst="139.133.210.32")/TCP()
2 >>> a=sr1(p)
3 Begin emission:
4 . Finished sending 1 packets.
5 *
6 Received 2 packets, got 1 answers, remaining 0 packets
7 >>> a
<IP  version=4 ihl=5 tos=0x0 len=44 id=0 flags=DF frag=0 ttl=47 proto=tcp checksum=0xa98d
  src=139.133.210.32 dst=172.22.152.130 options=[] |<TCP  sport=http dport=ftp_data
  seq=3081101820 ack=1 dataofs=6 reserved=0 flags=SA window=29200 checksum=0xe9c7 urgptr
  =0 options=[('MSS', 1452)] |<Padding  load=':v' |>>>
9 >>> a.summary()
'IP / TCP 139.133.210.32: http > 172.22.152.130: ftp_data SA / Padding'
```

Listing 10: Create and Send an IP/TCP Packet

Evil Bit

*The evil bit is a fictional **IPv4 packet header field** proposed in RFC 3514 [2], a humorous April Fools' Day RFC from 2003 authored by Steve Bellovin. The RFC recommended that the last remaining unused bit, the "Reserved Bit," in the IPv4 packet header be used to indicate whether a packet had been sent with malicious intent, thus making computer security engineering an easy problem — simply ignore any messages with the evil bit set and trust the rest.*

– Wikipedia

└ Evil Bit

The evil bit is a fictional IPv4 packet header field proposed in RFC 3514 [2], a humorous April Fools' Day RFC from 2003 authored by Steve Bellovin. The RFC recommended that the last remaining unused bit, the "Reserved Bit," in the IPv4 packet header be used to indicate whether a packet had been sent with malicious intent, thus making computer security engineering an easy problem — simply ignore any messages with the evil bit set and trust the rest.

— Wikipedia

If you enjoy the concept of the evil bit, you may like to also check out [5]: TCP Option to Denote Packet Mood. For example happy packets which are happy because they received their ACK return packet within less than 10ms. Or the Sad Packets which are sad because they faced retransmission rates greater than 20% of all packets sent in a session.

Evil Bit

Setting the Evil Bit with Scapy

- The flags in the IP header are just an attribute you can modify:

```
>>> i = IP()  
>>> i.flags = 'evil'
```

Listing 11: Setting the Evil Bit on an IPv4 Header with Scapy

PATHspider Plugins

Directory Layout

todo: Write this

PATHspider Plugins

ForgeSpider

```
class EvilBit(ForgeSpider, PluggableSpider):  
    name = "evilbit"  
    description = "Evil bit connectivity testing"  
    version = "0.0.0"  
    chains = [...]  
    connect_supported = [...]  
    packets = 2  
  
    def forge(self, job, seq):  
        ...
```

Listing 12: Outline for Evil Bit plugin using ForgeSpider

PATHspider Plugins

Forging the Packets

todo: Write this

PATHspider Plugins

Testing Packet Forging

todo: Write this

PATHspider Plugins

The Observer

todo: Write this

PATHspider Plugins

Library Observer Functions

todo: Write this

PATHspider Plugins

Observing the Evil Bit

todo: Write this

PATHspider Plugins

Putting it Together

todo: Write this

PATHspider Plugins

Adding a Flow Combiner

todo: Write this

PATHspider Plugins

Running it For Real

todo: Write this

Target Lists

Types of Targets

todo: Write this

Target Lists

Using *hellfire*

```
1 irl@z:~$ hellfire
Usage:
3  hellfire --topsites [--file=<filename>] [--output=<individual|array|oneeach>] [--type=<
    host|ns|mx>] [--canid=<canid address>]
    hellfire --cisco [--file=<filename>] [--output=<individual|array|oneeach>] [--type=<
    host|ns|mx>] [--canid=<canid address>]
5  hellfire --citizenlab [--country=<cc>|--file=<filename>] [--output=<individual|array|
    oneeach>] [--type=<host|ns|mx>] [--canid=<canid address>]
    hellfire --opendns [--list=<name>|--file=<filename>] [--output=<individual|array|
    oneeach>] [--type=<host|ns|mx>] [--canid=<canid address>]
7  hellfire --csv --file=<filename> [--output=<individual|array|oneeach>] [--type=<host|ns
    |mx>] [--canid=<canid address>]
    hellfire --txt --file=<filename> [--output=<individual|array|oneeach>] [--type=<host|ns
    |mx>] [--canid=<canid address>]
```

Listing 13: *hellfire*'s Usage Help

```
irl@z~$ hellfire --cisco
```

Listing 14: Start Resolving the Cisco Umbrella List

Target Lists

Using *canid* for Additional Information

todo: Write this

Target Lists

Resolving for Real

todo: Write this

Explicit Congestion Notification

Using the Built-In Plugin

todo: Write this

Simple Analysis with PATHspider

Up next:

Path Transparency Observatory (PTO)

Don't delete your PATHspider results, you'll need them in the next session.

References I

- [1] J. Arkko, M. Cotton, and L. Vegoda.
IPv4 Address Blocks Reserved for Documentation.
RFC 5737 (Informational), January 2010.
- [2] S. Bellovin.
The Security Flag in the IPv4 Header.
RFC 3514 (Informational), April 2003.
- [3] B. Claise, B. Trammell, and P. Aitken.
Specification of the IP Flow Information Export (IPFIX) Protocol for
the Exchange of Flow Information.
RFC 7011 (Internet Standard), September 2013.

References II

- [4] Ana Custura, Gorrry Fairhurst, and Iain Learmonth.
Exploring usable path MTU in the Internet.
In Proceedings of the 2018 Network Traffic Measurement and Analysis Conference (TMA '18). IEEE, Jun 2018.
- [5] R. Hay and W. Turkal.
TCP Option to Denote Packet Mood.
RFC 5841 (Informational), April 2010.
- [6] Mirja Kühlewind, Sebastian Neuner, and Brian Trammell.
On the state of ECN and TCP options on the Internet.
In Proceedings of the Passive and Active Measurement Conference,
pages 135–144, Hong Kong, China, 2013.

References III

- [7] Iain R. Learmonth, Brian Trammell, Mirja Kühlewind, and Gorrry Fairhurst.
PATHspider: A tool for active measurement of path transparency.
In *Proceedings of the 2016 Applied Networking Research Workshop*,
pages 62–64, July 2016.
- [8] Brian Trammell, Mirja Kühlewind, Damiano Boppart, Iain Learmonth,
Gorrry Fairhurst, and Richard Scheffenegger.
Enabling internet-wide deployment of explicit congestion notification.
In *Proceedings of the Passive and Active Measurement Conference*,
pages 193–205, New York, USA, 2015.